

# Information Retrieval

Joseph Fitzpatrick

13/10/18

## 1 Introduction

The following document will outline the process by which I created a Search Engine on the Cranfield Index using lucene. The results were evaluated using the "Trec Eval" library to test for a wide variety of scores.

## 2 Implementation

### 2.1 Preprocessing Documents

The first step in creating the search engine was to process and index all of the documents required. For this I created a simple parser for the entire "Cranfield 1400" file. The parser would read the file in line by line. Each line would then be added to an variable depending on it's context, ie. if it was a Title, Author, Bibliography or Content. Once the context looped around to the starter context, the variables would be used to create a new document.

Once all the documents were created, they were added to an 'IndexWriter' object and then written to a 'Directory' stored on Disk. The directory was then analysed using a custom analyser that I created. The analyser removed stop words, normalised characters to lower case and stemmed the words.

### 2.2 Searching the Index

The next step was to begin searching the index. As the index contained a number of fields for each document which need to be searched I used the "MultiFieldQueryParser". This query parser allows you to search across all of the fields at once. It does however treat all of the fields as equal, giving each an equal weighting. I decided to change this and alter the weighting for each field.

Field	Weighting
Index	0
Title	0.35
Bibliography	0.01
Author	0.03
Content	0.61

These score were calculated by performing some small calculations to identify the optimal weightings to give the best results. Rarely do people search for the author of a book or it's bibliography information and therefore, both of those received very low weightings. Giving the higher weightings for the content and title was based on the fact that people search for those more often.

The Search queries were parsed from the "Cran.qry" file. They were then trimmed and parsed using the custom analyser to remove stop words and stem the words. Once the queries were in the correct format for searching they were then fed into the IndexSearcher Object. This searched the index for the queries and returned the top 30 results identified.

## 2.3 Ranking Results

The results that are returned from the IndexSearcher object are ranked based on their relevancy score. This score indicates how relevant a Document is based on the query. The documents are ranked highest to lowest.

## 2.4 Custom Analyser

Once I had the initial minimal search engine working I began looking at ways to improve the results I obtained. To do this I used the "Luke" index visualisation tool to view the directory created. With this I could see that the Standard Analyser was lacking in a few areas. With some research on lucene I began to create the analyser by extending the 'StopwordAnalyzerBase' class. This class allowed me to customize the stop-words I wanted to exclude, while also allowing me to extend the analysers functionality. I decided to use the 'English Stop Words Set' as it contained the majority of the words 'Luke' was showing. I also added a lower-case normaliser to the analyser, as well as a stemming filter to stem words.

## 3 Results

Providing results for the effectiveness of the search engine was carried out using "Trec Eval". This tool works by performing some calculations against your results and the optimal results. I also compared different scoring approaches in my model.

Scoring	MAP	Recall	Reciprocal Rank	P5
BM25	0.412	0.8320	0.832	0.4587
Boolean	0.3013	0.7257	0.7056	0.3556
Classic/VSM	0.3805	0.8192	0.8024	0.4213

From these results, it's clear that using BM25 similarity scoring produces the best results across the board, and boolean similarity produces the worst results. BM25 is an improved version of the Classic Similarity TF\*IDF scoring mechanism usually used within lucene. It uses a probabilistic score for determining if a

user will find a certain document relevant when scoring the documents. Boolean Similarity gives the worst results as it only scores a document if it fully matches the query. Therefore no partial results are returned leaving gaps in the information retrieved. The Vector Space Model (VSM) or Classic scoring performed well, however it wasn't as good as the more advanced similarity method which did more than just calculating the TF\*IDF.

## 4 Conclusion

From writing the Search Engine with the lucene library, I have been able to explore the information retrieval process in greater detail. Being able to compare the impact of different techniques, such as stemming and stop-word removal, showed their significance. As a whole, it take a combination of a lot of different factors to create a search engine capable of returning significant results.